



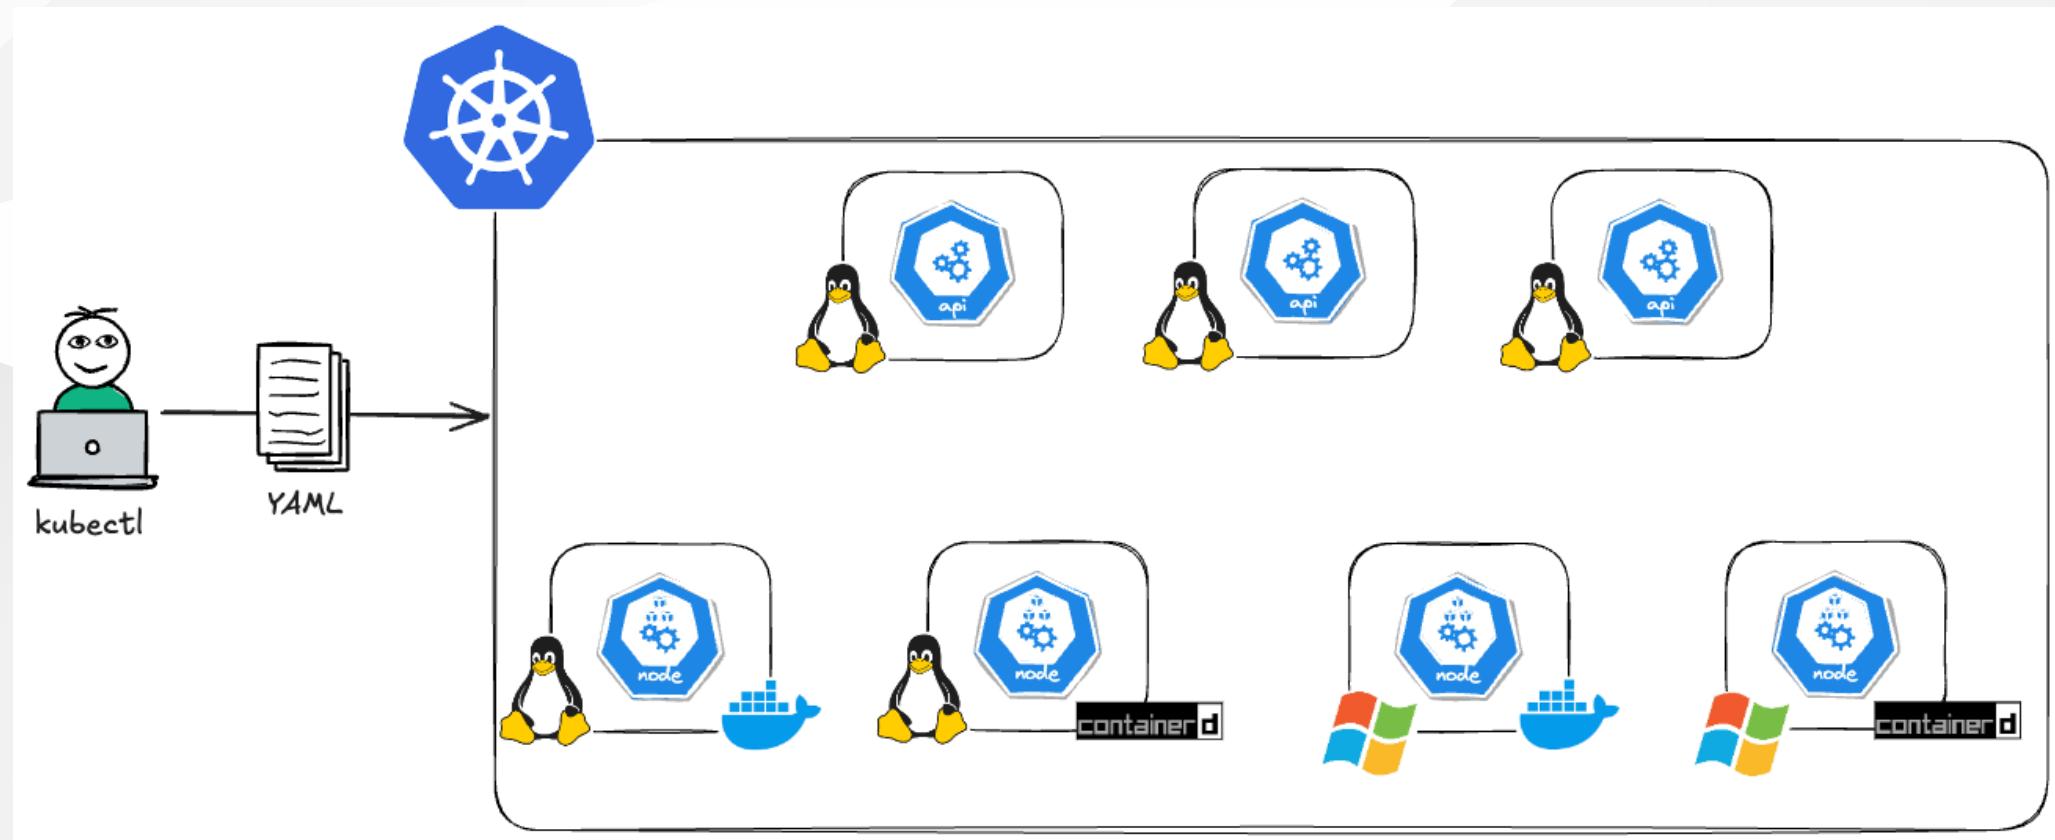
k8s Workshop

Part 1
Running containers with Pods and
Deployments

Why k8s is so popular ?



API and Cluster

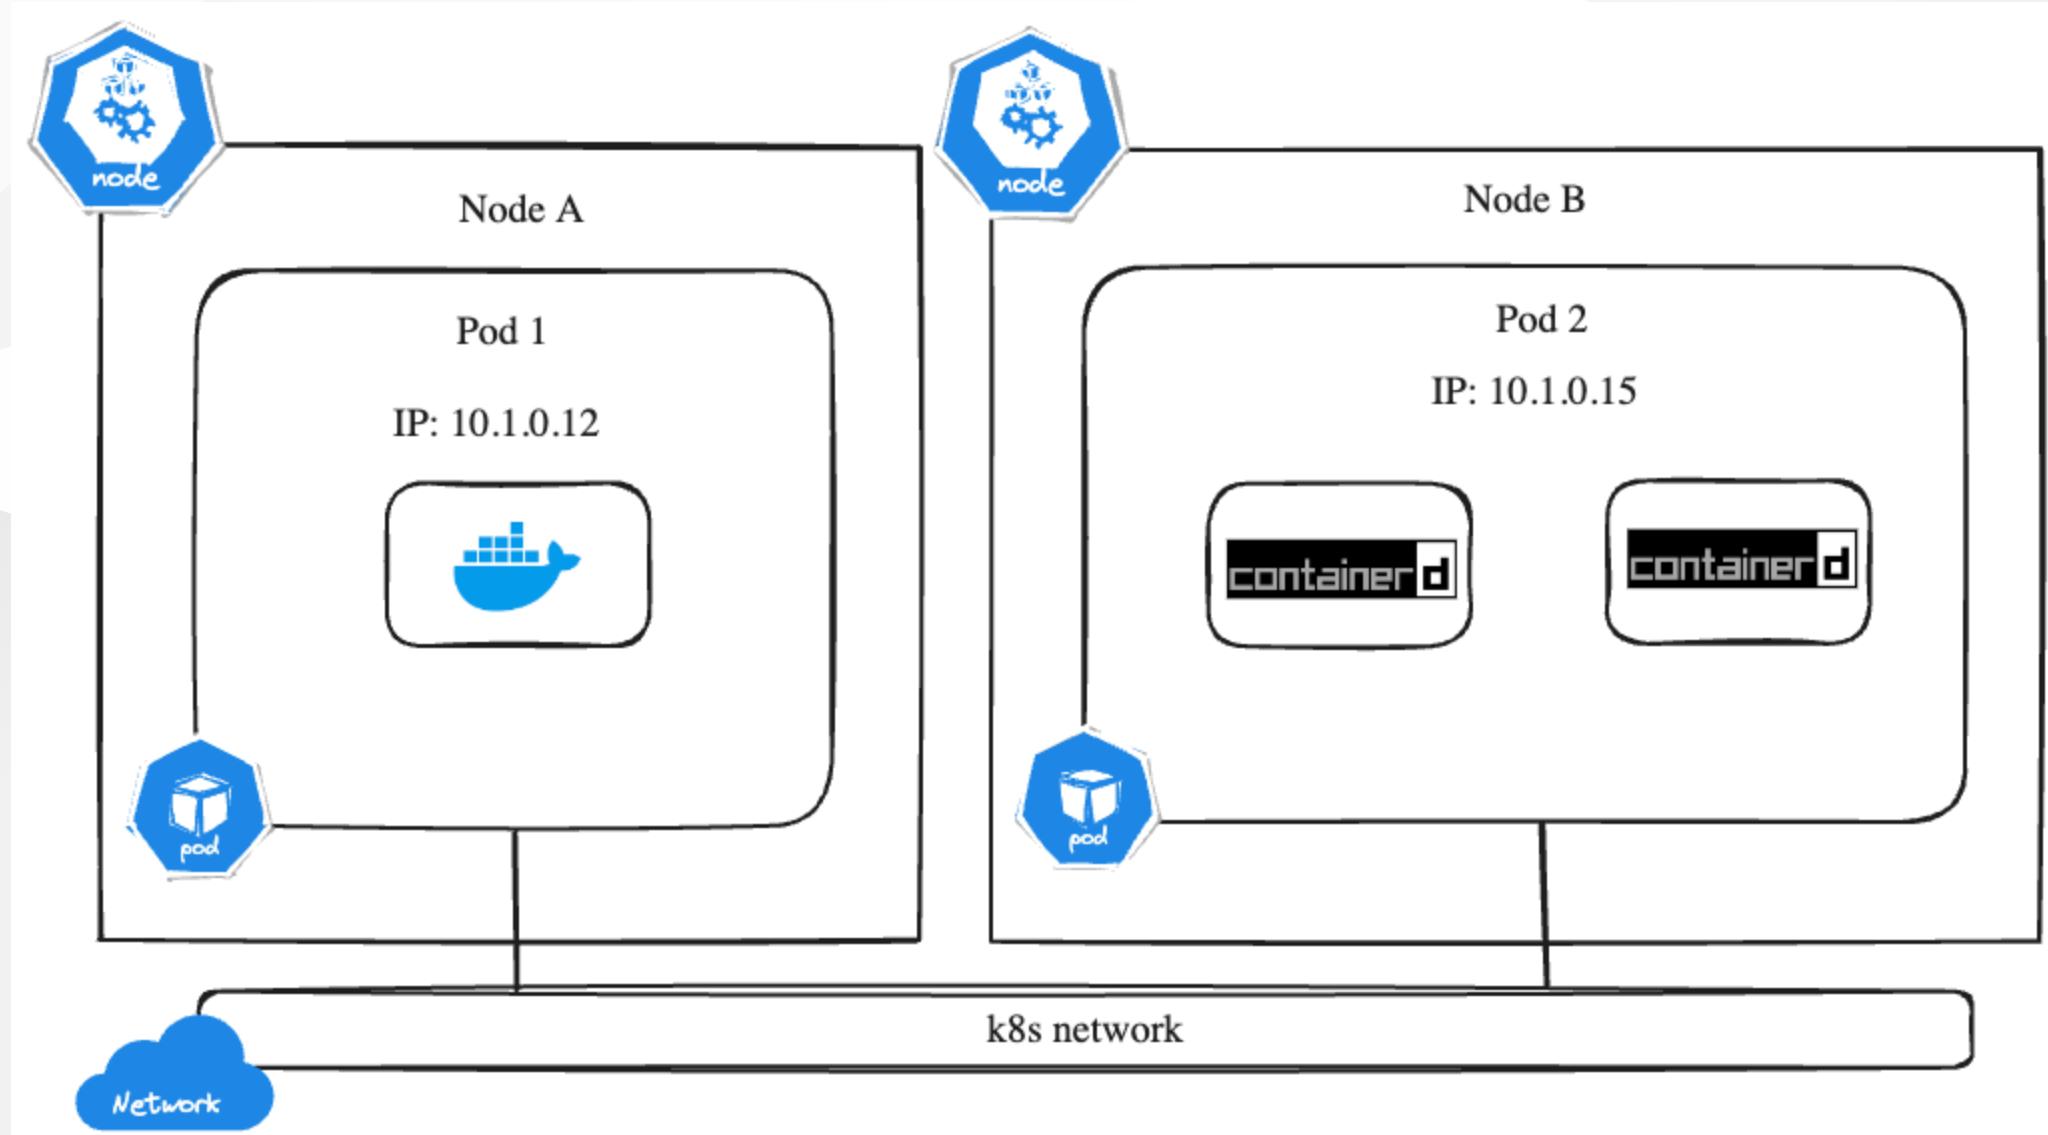


Containers & Pods

- Container is virtualized environment
- Pod wraps one or more containers
- Pod is a unit of compute
- Pod has a unique virtual IP
- Pods communicate over the k8s network
- Containers in a Pod share the same network address

Why a Pod is the smallest deployable unit in k8s ?

- k8s runs Pods not Containers
- Containers are managed by the container runtime of the node (Docker, containerd, etc.)
- Nodes manages containers using the CRI API
- The CRI API is standard for all container runtimes

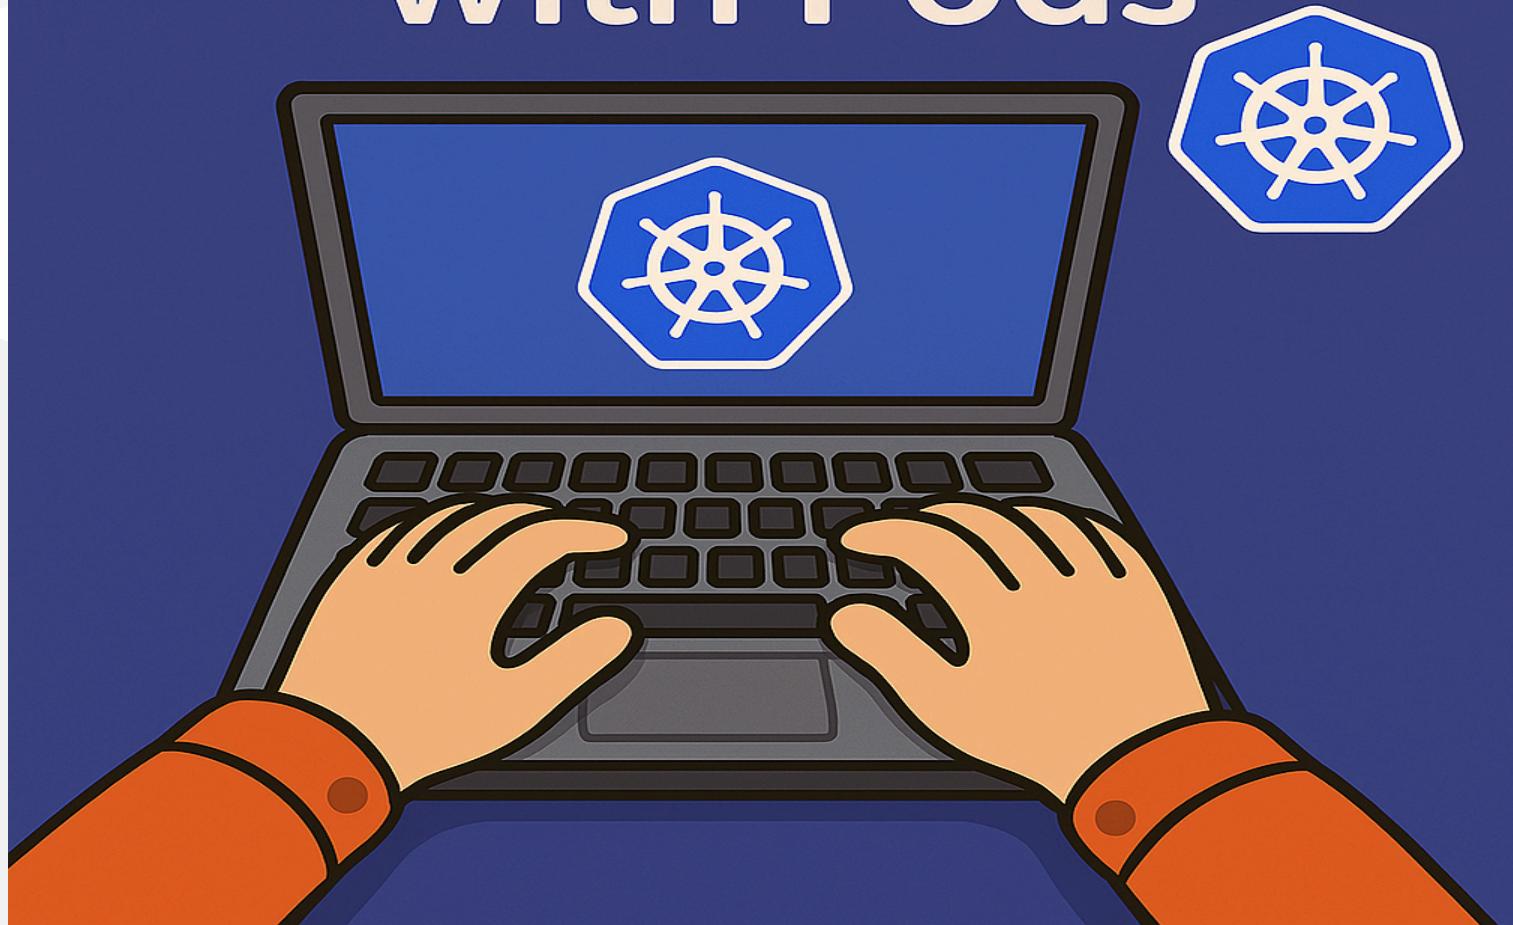


kubectl

kubectl is cli tool for communicating with the k8s cluster's control plane using the k8s API. Some of the main kubectl commands:

- get: get information about resources
- describe: get detailed information about a resource
- run: run a resource
- create: create a resource
- apply: create or update a resource
- delete: delete a resource

HANDS-ON with Pods



Let's create a Pod

```
kubectl run hello-world-1 --image=salvovitale/k8s-ws-app-hello-world:latest --restart=Never
```

check the pod

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-world-1	1/1	Running	0	39s

Use custom-columns

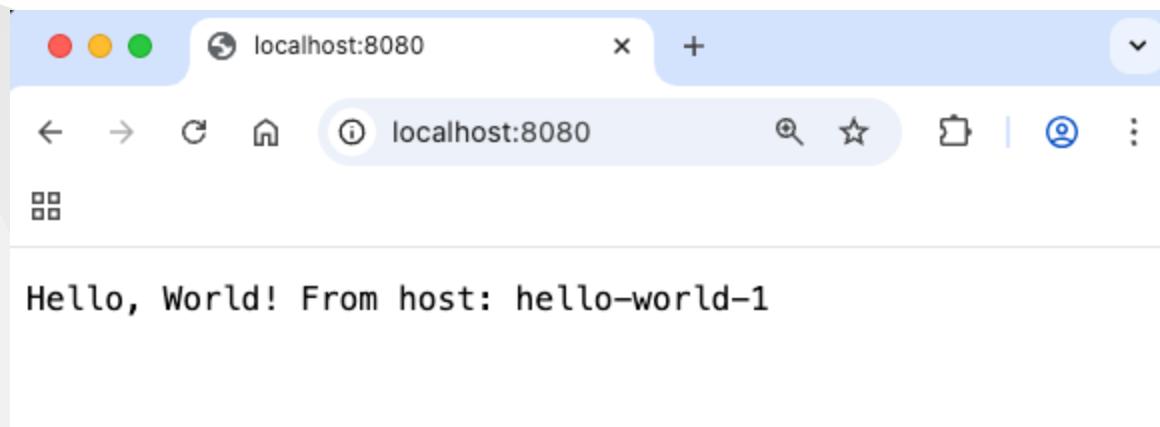
```
kubectl get pod hello-world-1 --output \
custom-columns=NAME:metadata.name,NODE_IP:status.hostIP,POD_IP:status.podIP
NAME          NODE_IP        POD_IP
hello-world-1  172.19.0.3   10.244.1.4
```

Search for specific information

```
kubectl get pod hello-world-1 -o jsonpath='{.status.containerStatuses[0].containerID}'
containerd://b4c566898b5e875a95cd475a71126da585e6f99afebb79f51dd88ca62eb5e53c%
```

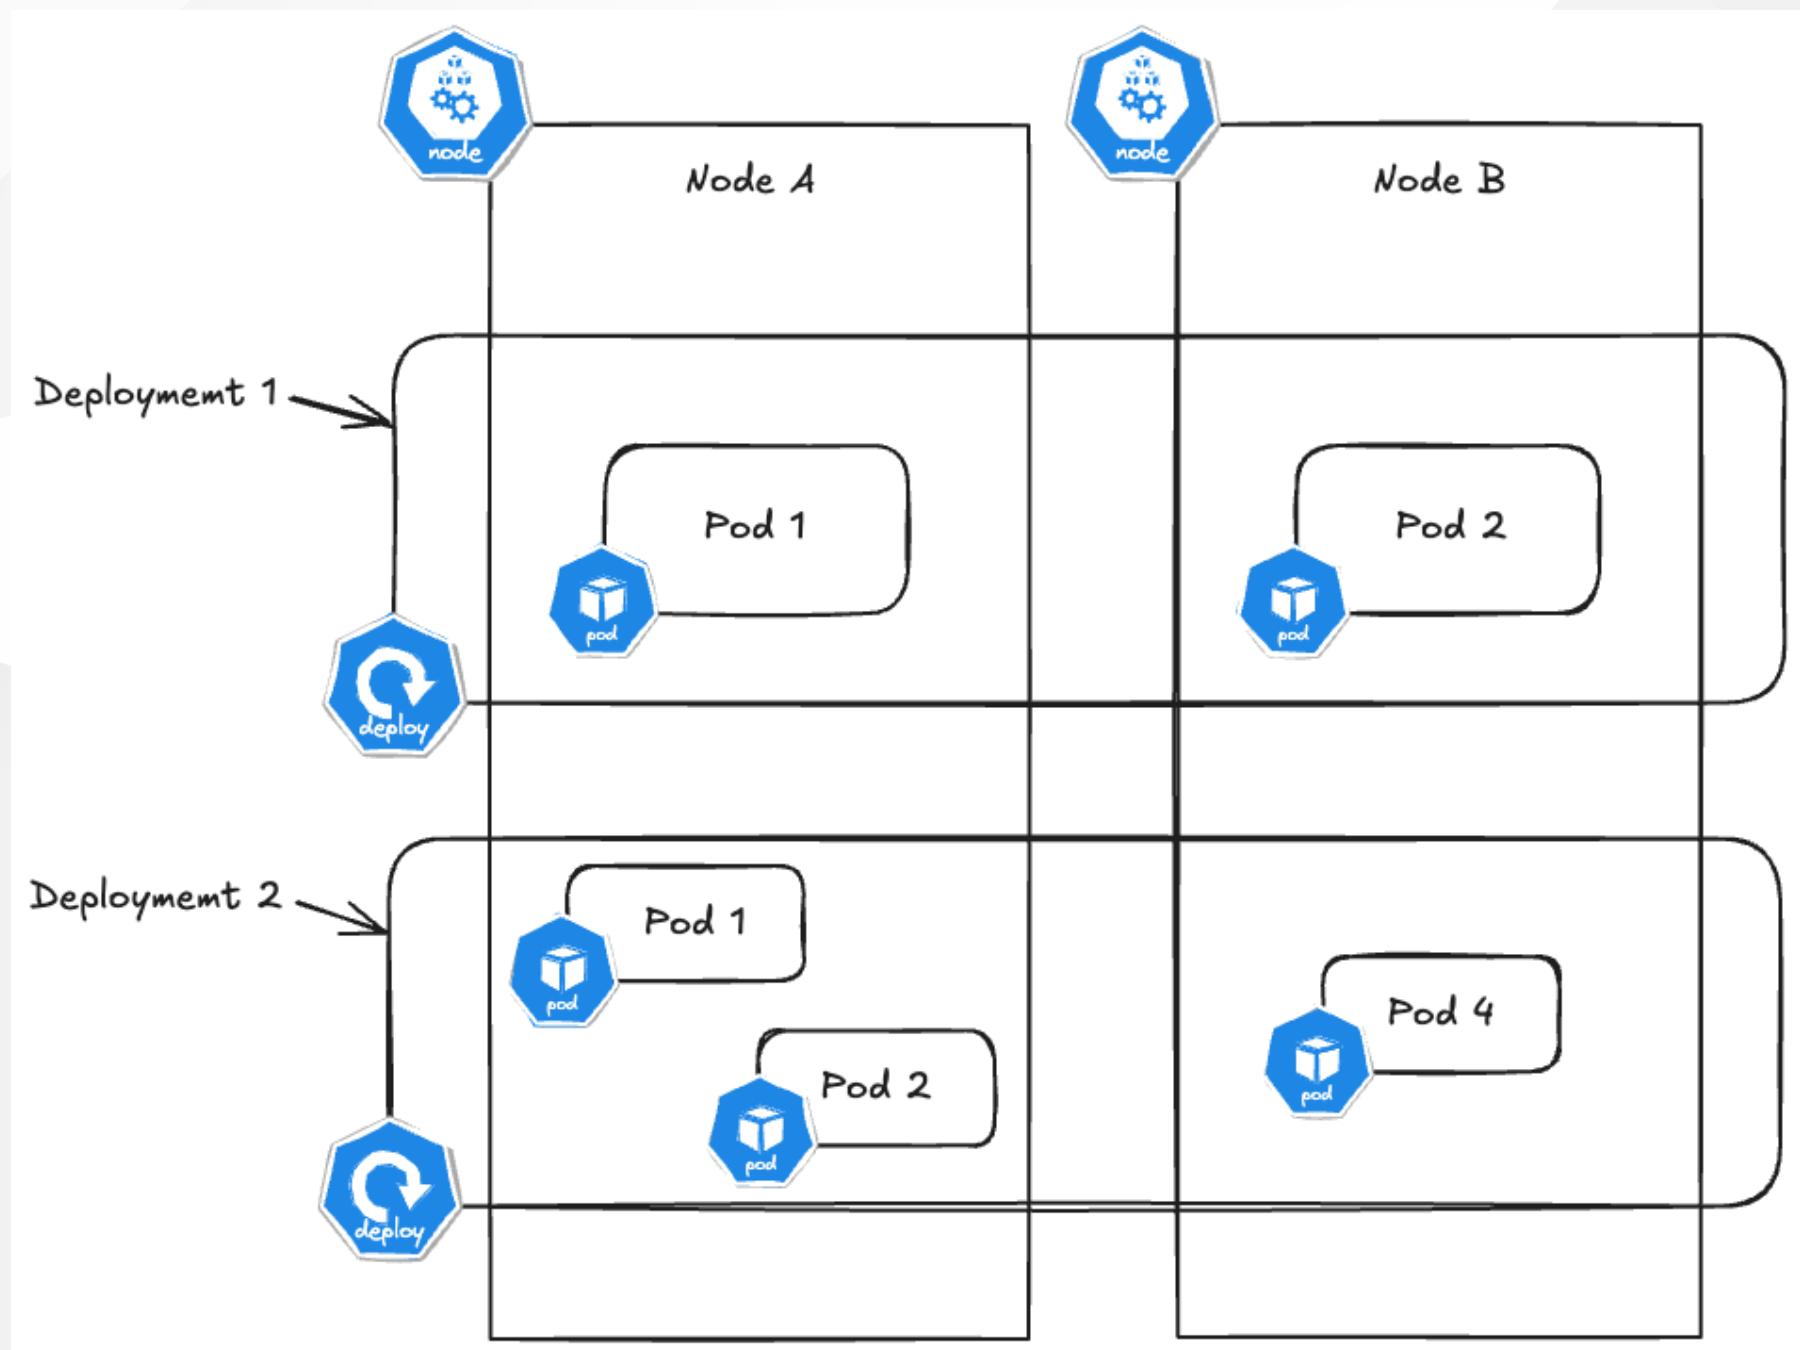
Port-forward

```
kubectl port-forward pod/hello-world-1 8080:8080  
Forwarding from 127.0.0.1:8080 -> 8080  
Forwarding from [::]:8080 -> 8080
```



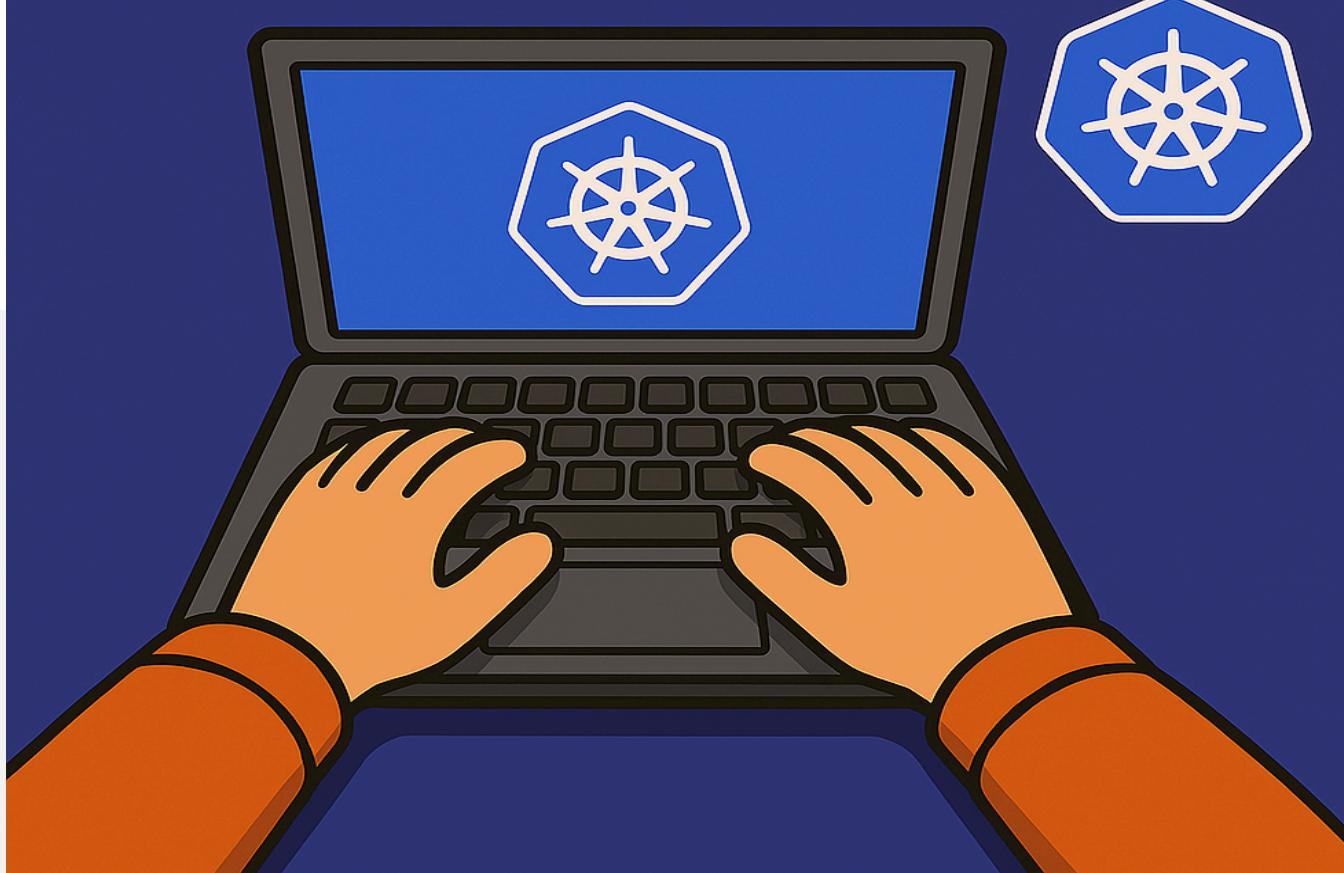
Deployments

- Pods are too simple
- Deployments is a controller that manages Pods
- Controller ensures that its managed resources are running and healthy
- Controller uses the k8s API to watch the state of its managed resources



HANDS-ON

with Deployments



Create a deployment

```
kubectl create deployment hello-world-2 --image=salvovitale/k8s-ws-app-hello-world:latest
```

check the deployment

```
kubectl get deploy hello-world-2
```

```
kubectl describe deploy hello-world-2
```

Resources Relationship and Labels

- k8s resources can have labels
- Labels are key-value pairs attached to k8s resources
- k8s uses labels to loosely couple resources (Deployments <--> Pods)
- controllers use labels to identify the resources they manage
- controllers don't need to maintain the list of resources they manage

Get the labels for the deployment

```
kubectl get deploy hello-world-2 -o jsonpath='{.spec.template.metadata.labels}'
```

we can use the labels to identify the pods

```
kubectl get pod -l app=hello-world-2
```

lets inspect the current labels situation in pods

```
kubectl get pods -o custom-columns=NAME:metadata.name,LABELS:metadata.labels
NAME                               LABELS
hello-world-1                      map[run:hello-world-1]
hello-world-2-69c988857-bdj5w      map[app:hello-world-2 pod-template-hash:69c988857]
```

lets override the label of a pod linked to the deployment

```
kubectl label pods -l app=hello-world-2 --overwrite app=hello-world-x
```

lets inspect the new labels situation in pods

```
kubectl get pods -o custom-columns=NAME:metadata.name,LABELS:metadata.labels
NAME                               LABELS
hello-world-1                      map[run:hello-world-1]
hello-world-2-69c988857-9fkxf      map[app:hello-world-x pod-template-hash:69c988857]
hello-world-2-69c988857-bdj5w      map[app:hello-world-2 pod-template-hash:69c988857]
```

lets revert back the situation

```
kubectl label pods -l app=hello-world-x --overwrite app=hello-world-2
```

Manifests

- complete description of your app
- written in JSON or YAML
- versioned and tracked in source control
- create the same deployment on any k8s cluster
- declarative not imperative
- `kubectl apply -f <manifest>`

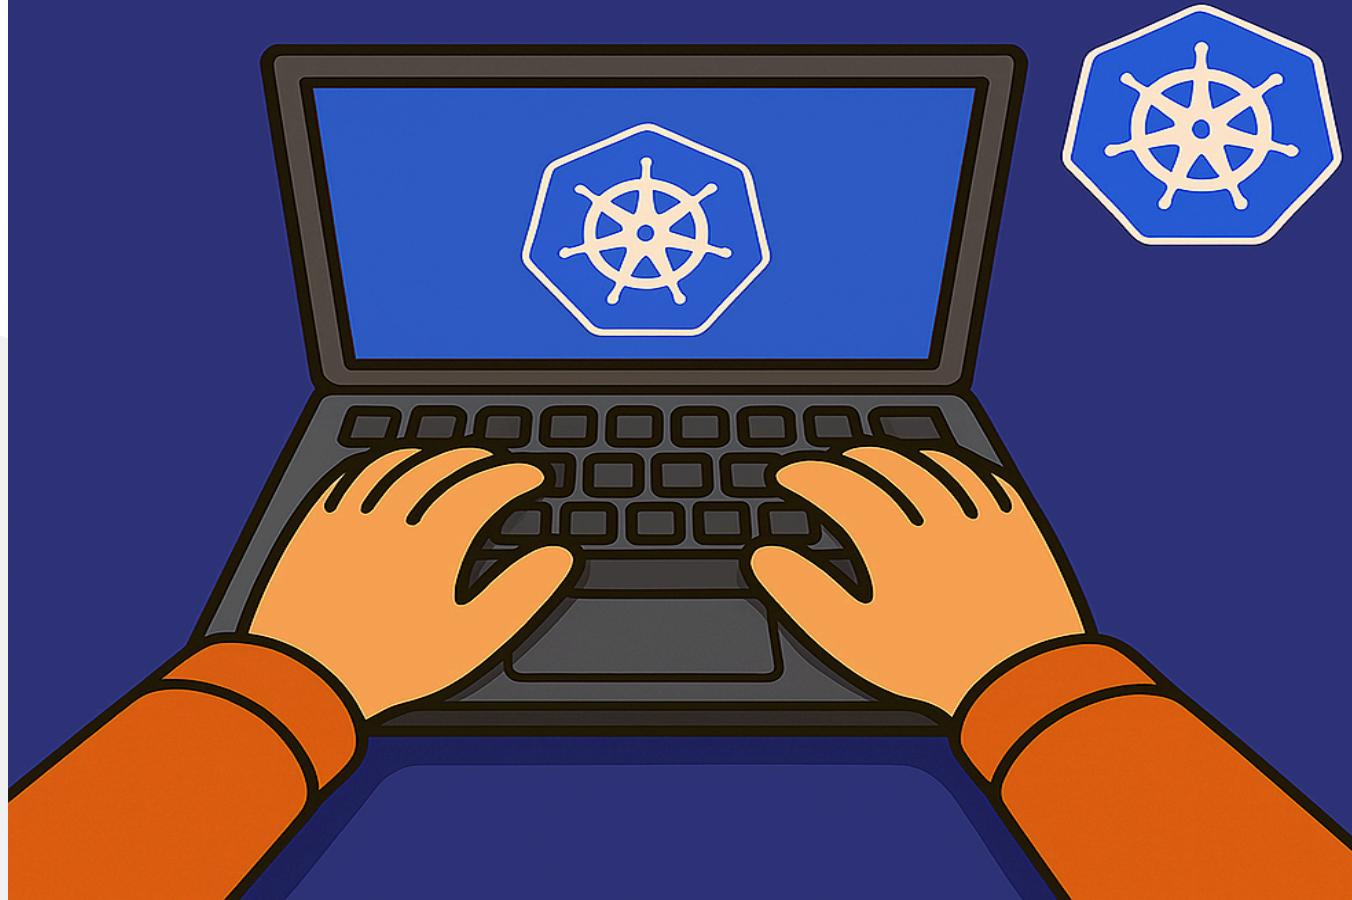
Pod manifest

```
# Manifests always specify the version of the Kubernetes API and the type of resource.
apiVersion: v1
kind: Pod
# Metadata for the resource includes the name (mandatory) and labels (optional).
metadata:
  creationTimestamp: null
  labels:
    run: hello-world-3
    name: hello-world-3
# The spec is the actual specification for the resource.
# For a Pod the minimum is the container(s) to run,
# with the container name and image.
spec:
  containers:
  - image: app-hello-world:latest
    name: hello-world-3
```

Deployment manifest

```
# Deployments are part of the apps version 1 API spec.
apiVersion: apps/v1
kind: Deployment
# Metadata for the resource includes the name (mandatory) and labels (optional).
metadata:
  labels:
    app: hello-world-4
    name: hello-world-4
# The spec includes the label selector the Deployment uses it to find its own managed resources. I'm using the app label,
# but this could be any combination of key-value pairs.
spec:
  selector:
    matchLabels:
      app: hello-world-4
# The template is used when the Deployment creates a Pod
  template:
# Pods in a Deployment don't have a name, but they need to specify labels that match the selector in the spec.
    metadata:
      labels:
        app: hello-world-4
# The Pod spec lists the container name and image
    spec:
      containers:
        - image: app-hello-world:latest
          name: app-hello-world
```

HANDS-ON with Manifests



Lets use the cli to generate a manifest for the pod we created before

```
kubectl run hello-world-3 \
--image=salvovitale/k8s-ws-app-hello-world:latest \
--image-pull-policy=IfNotPresent \
-o yaml \
--dry-run=client > pod.yaml
```

and apply it

```
kubectl apply -f pod.yaml
```

lets now generate the manifest for a deployment

```
kubectl create deployment hello-world-4 \
--image="app-hello-world:latest" \
-o yaml \
--dry-run=client > deployment.yaml
```

and apply it

```
kubectl apply -f deployment.yaml
```

Interactive shell

We can run an interactive shell in the pod container

```
kubectl exec -it hello-world-3 -- sh
```

we can also target a deployment

```
kubectl exec -it deploy/hello-world-4 -- sh
```

Inspect Logs

We can check the logs of the application running in the pod

```
kubectl logs pod/hello-world-3 -f
```

again also in this case we can refer directly to the deployment

```
kubectl logs deploy/hello-world-4 -f
```

Filesystem

Interacting with the filesystem of pod

```
kubectl cp temp/text.txt hello-world-3:/app/text.txt
```

also works the other way around to copy a file from the pod to the local filesystem

Delete Resources

Lets delete all the pods

```
kubectl delete pods --all
```

The pod managed by the deployment will be re-created. To delete that resource we need to delete the controlling resource. In this case the deployment

```
kubectl delete deploy --all
```

Some final info

- ohmyzsh kubectl alias

```
alias | grep "kubectl"
```

Lab

- Use the "kubectl run" to create a pod with the official Nginx image
- Port-forward to the pod
- Use the "kubectl create" to create a deployment with the official Nginx image
- Use kubectl to generate basic manifests for both pod and deployment
- Interact with the filesystem to modify/replace the nginx landing page
- Repeat with an hello-world app in your technology of choice